

Syed Abdul Muqet Mujeeb – AI Engineer

a.muqetmujeeb@gmail.com | +91 9010830602 | github.com/MuqetMujeeb | linkedin.com/in/muqetmujeeb | Hyderabad

Personal Summary

AI Engineer specializing in production voice AI agents, RAG systems, and real-time backend pipelines. Experience spans LLM integration, vector search, REST/WebSocket APIs, and containerized cloud deployments on AWS, with a focus on reliability and maintainable architecture.

Professional Experience

AI Engineer

July 2026 – Present

Crypto Association Georgia

Hyderabad, India (Remote)

- Built **LioraAI** (lioraai.net), an AI-powered CRM that autonomously places outbound calls, qualifies leads, and routes them to partner call centers — Python backend orchestrating **LiveKit** real-time media over a **SIP** telephony provider.
- Engineered the voice-agent pipeline (**Gemini** reasoning, **OpenAI STT/TTS**), tuning barge-in, endpointing, and streaming buffers to sustain **~800ms** turn latency on live PSTN calls.
- Automated **disposition classification** (call complete, callback requested, voicemail) via LLM-based post-call analysis, exporting filtered, call-center-ready lead batches as CSV.
- Shipped call history, callback scheduling, per-call **cost analytics**, and role-based **agent/admin portals**; deployed on **AWS EC2** with **auto-scaling** behind a load balancer.
- Built a real-time face- and voice-alteration app (**Deep-Live-Cam**, **w-okada RVC**) streamed over **WebRTC** from GPU-backed AWS instances.

AI Developer

December 2025 – July 2026

SMARTnCODE Technologies

Hyderabad, India (On-site)

- Built a full-stack AI legal chatbot with a **RAG** pipeline (Pinecone + BGE-Large + Gemini API), achieving **85% response accuracy** and **<3s latency**; a rolling 10-message summarization strategy cut token consumption **~40%**.
- Delivered multi-modal input — OCR document upload and **Whisper** speech-to-text — via a **FastAPI** backend with Redis session management and SQL Server storage.
- Developed a real-time conversational avatar (SoulX-Flashhead on RunPod cloud GPU, WebRTC) at **~2s end-to-end latency**; integrated Gemini STS, resolved audio-visual sync conflicts, and designed a scalable multi-user session and database layer.
- Contributed to R&D on an AI video generation platform, integrating diffusion models (Google Veo, Kling AI, Seedance) into a unified cinematic content pipeline.

Project Experience

MindCanvas (Supabase, LangChain, Cytoscape.js, FastAPI, React, OpenAI)

- Built an AI knowledge system converting browsing data into clustered knowledge graphs paired with a **RAG-based learning assistant**, with interactive nodes to explore details and relationships, enhancing learning retention.
- Achieved 3x faster retrieval and a **75% improvement** in identifying latent relationships across datasets.
- Developed a full-stack application and Chrome extension for AI-powered content analysis and visualization.

HireSense (Next.js, Supabase, PostgreSQL, GPT-4, Groq API, ElevenLabs)

- Built an AI hiring platform that automated cross-validation of CVs, LinkedIn, and GitHub data, reducing candidate verification time by **60%**.
- Designed a GPT-4-powered Certainty Score using sentiment analysis to quantify candidate credibility; integrated AI-driven empathetic interviews via ElevenLabs and Groq for behavioral insights.

Education

B.E. Computer Science (Artificial Intelligence & Machine Learning)

GPA: 8.5

Lords Institute of Engineering & Technology

Nov 2022 – July 2026

Skills

Languages/Backend	Python, SQL, React, Next.js, Git, FastAPI, REST APIs
AI/ML	LLMs, RAG, LangChain, OpenCV, Transformers, PyTorch, TensorFlow, Scikit-learn, Pandas
Infra/Data	Docker, AWS (EC2, Auto Scaling), RunPod, PostgreSQL, SQL Server, Redis, Supabase, Qdrant, Pinecone
Voice/Realtime	LiveKit, WebRTC, WebSockets, SIP/PSTN, STT/TTS, Whisper, Gemini STS

Achievements

- Competed in **6+ hackathons**, building AI-driven solutions under real-world time constraints.
- Active LeetCode problem solver — consistently practising data structures and algorithms.